

# ADL Final Project Presentation

Qi Feng (qf2140)



# Goal

- Given a gigapixel pathology image (slide), the goal is to classify if the image contains tumor and localize the tumors for a pathologist's review.



# Preprocessing

- Data loading
- Patch extraction
- Train-test-split
- Class Imbalance



# Data Loading

- Sanity check
  - Number of slides == Number of tumor masks
    - Missing mask for ['/content/drive/My Drive/adl/slides/tumor\_038.tif']
    - Remedy: drop slide/mask with no match.
    - Total slides: 21
  - Slide dimensions == Tumor mask dimensions
    - Minimum dimensions: 8 (level 0-7)



# Patch Extraction

- Parameters used for sliding-window
  - LEVELS = [3, 4]
  - PATCH\_SIZE = 300
  - STRIDE = 128
- Train-test-split
  - Randomly select two slides for testing: ['tumor\_075', 'tumor\_084']
  - Draw **overlapping** patches for the **training** set (stride = 128).
  - Draw **non-overlapping** patches for the **test** set (stride = 300).

# Method

- For all levels, iterate the center (x, y) of the patches
  - Coordinates are multiplied by scales
- Pad at higher levels when (x, y) is out of border (i.e., the maximum dimensions of the lowest level)
- Skip the non-tissue areas
  - Threshold = 40%
  - Record the indices to take labels/skip patches in the next level accordingly

```
def _extract_patch(slide_img, max_x, max_y, check, dropped,
                  scale=1, patch_size=PATCH_SIZE, stride=STRIDE):
    """
    Helper function to extract patch at a level.
    """
    max_x_cur = max_x * scale # Max dim for current level
    max_y_cur = max_y * scale
    # stride *= scale
    half_patch = patch_size // 2
    patches = []

    idx = 0 # used to track dropped patch index
    # Iterate the center of the patch
    for y in enumerate(np.arange(half_patch, max_y - half_patch, stride)):
        y *= scale
        for x in enumerate(np.arange(half_patch, max_x - half_patch, stride)):
            x *= scale

            if not check:
                if idx in dropped:
                    idx += 1
                    continue
                else:
                    idx += 1

            # Padding will only be executed at higher levels
            # Pad when out of border
            y_top = int(max(0, y - half_patch))
            y_bottom = int(min(max_y_cur, y + half_patch))
            x_left = int(max(0, x - half_patch))
            x_right = int(min(max_x_cur, x + half_patch))

            patch = slide_img[y_top : y_bottom, x_left : x_right, :]

            if patch.shape[0] != patch_size:
                if y_top == 0:
                    patch = pad_along_axis(patch, half_patch - y, 0)
                if y_bottom == max_y_cur:
                    patch = pad_along_axis(patch, half_patch - (max_y_cur - y), 0)
            if patch.shape[1] != patch_size:
                if x_left == 0:
                    patch = pad_along_axis(patch, half_patch - x, 1)
                if x_right == max_x_cur:
                    patch = pad_along_axis(patch, half_patch - (max_x_cur - x), 1)

            if check: # if at min level, check and skip non-tissue areas
                tissue_pixels = find_tissue_pixels(patch)
                percent_tissue = len(tissue_pixels) / float(patch.shape[0] * patch.shape[1]) * 100
                if percent_tissue < 40:
                    dropped.append(idx)
                    idx += 1
                    continue
                idx += 1

            assert patch.shape == (PATCH_SIZE, PATCH_SIZE, 3)
            patches.append(patch)

    return patches, dropped
```

# Labelling

- Draw patches correspondingly on the mask image
- For each input patch, we predict the label of the center  $128 \times 128$  region.
- We label a patch as **tumor ('1')** if at least one pixel in the center region is annotated as tumor
  - Otherwise, **Non-tumor ('0')**

```
def crop_center(img, crop_x=128, crop_y=128):  
    """  
    Crop the center area for an image.  
    """  
    y, x, _ = img.shape  
    start_x = x // 2 - (crop_x // 2)  
    start_y = y // 2 - (crop_y // 2)  
    return img[start_y:start_y + crop_y, start_x:start_x + crop_x, :]
```

```
def _extract_patch_labels(mask_img, max_x, max_y, dropped,  
                          scale=1, patch_size=PATCH_SIZE, stride=STRIDE):  
    """  
    Helper function to extract patch labels at a level.  
    """  
    max_x_cur = max_x * scale # Max dim for current level  
    max_y_cur = max_y * scale  
    # stride *= scale  
    half_patch = patch_size // 2  
  
    labels = []  
  
    idx = 0 # used to track dropped patch index  
    for _, y in enumerate(np.arange(half_patch, max_y - half_patch, stride)):  
        y *= scale  
        for _, x in enumerate(np.arange(half_patch, max_x - half_patch, stride)):  
            x *= scale  
            if idx in dropped:  
                idx += 1  
                continue  
            else:  
                idx += 1  
  
            # Padding will only be executed at higher levels  
            # Pad when out of border  
            y_top = int(max(0, y - half_patch))  
            y_bottom = int(min(max_y, y + half_patch))  
            x_left = int(max(0, x - half_patch))  
            x_right = int(min(max_x, x + half_patch))  
  
            patch = mask_img[y_top : y_bottom, x_left : x_right, :]  
  
            if patch.shape[0] != patch_size:  
                if y_top == 0:  
                    patch = pad_along_axis(patch, half_patch - y, 0)  
                if y_bottom == max_y_cur:  
                    patch = pad_along_axis(patch, half_patch - (max_y_cur - y), 0)  
            if patch.shape[1] != patch_size:  
                if x_left == 0:  
                    patch = pad_along_axis(patch, half_patch - x, 1)  
                if x_right == max_x_cur:  
                    patch = pad_along_axis(patch, half_patch - (max_x_cur - x), 1)  
  
            # Crop the center 128 region to label  
            patch = crop_center(patch)  
  
            lab = 0 if np.all(patch == 0) else 1  
            labels.append(lab)  
  
    return labels
```



# Loop Over Levels At All Slides

For all slides/tumor masks

21 Slides

21 Tumor masks

For all levels

## Level 3


Extract patches (filter out non-tissue, return indices)  
→ extract labels (center 128 region)

## Level 4


Extract (context) patches (matching level 3, pad if out of border)  
→ extract labels (center 128 region)



Level: 3, Label: 0




Level: 4, Label: 0




Level: 3, Label: 0




Level: 4, Label: 0




Level: 3, Label: 0




Level: 3, Label: 0



Level: 4, Label: 0




Level: 3, Label: 0



Level: 4, Label: 0



Level: 3, Label: 0



Neighboring training patches (overlapped)

Neighboring training patches, padded  
(overlapped)

Neighboring  
test patches,  
(non-overlapped)

```
def get_n_patch(slide_path, tumor_mask_path, level):
    slide = open_slide(slide_path)
    tumor_mask = open_slide(tumor_mask_path)

    # number of patches per col
    n_patch_h = slide.level_dimensions[level][0] // PATCH_SIZE
    # number of patches per row
    n_patch_v = slide.level_dimensions[level][1] // PATCH_SIZE


    return (n_patch_h, n_patch_v)
```

```
import matplotlib.image as mpimg

def concat_patches(patch_path, n_patch_h, n_patch_v, level):
    # Convert path into np array
    img = [mpimg.imread(i) for i in patch_path[level]]


    # 1. Stack patches horizontally
    partition = np.array_split(img, n_patch_v)
    rows = [np.hstack(p) for p in partition]
    # 2. Stack rows of patches vertically
    res = np.vstack(rows)

    return res
```



Sanity checks for patch extraction, tumor\_075.tif

# Train-test-split



# Class Imbalance

- Tumor patches are rare
- Solution: class weights

Class Counts



```
from sklearn.utils import class_weight

class_weights = class_weight.compute_class_weight('balanced',
                                                  [0, 1],
                                                  all_y_train[3])

class_weights = dict(enumerate(class_weights))
class_weights

{0: 0.6024266936299292, 1: 2.9407699901283317}
```



# Modeling

- Single scale, simple CNN
- Single scale, one Inception tower
- Multi-scales, two CNNs

# Single Scale, Simple CNN

```
def make_model():
    model = Sequential()
    model.add(Conv2D(32, (3, 3), activation='relu',
                    input_shape=(299, 299, 3)))
    model.add(MaxPooling2D())

    model.add(Conv2D(32, (3, 3), activation='relu'))
    model.add(MaxPooling2D())

    model.add(Dropout(0.5))

    model.add(Conv2D(64, (3, 3), activation='relu'))
    model.add(MaxPooling2D())

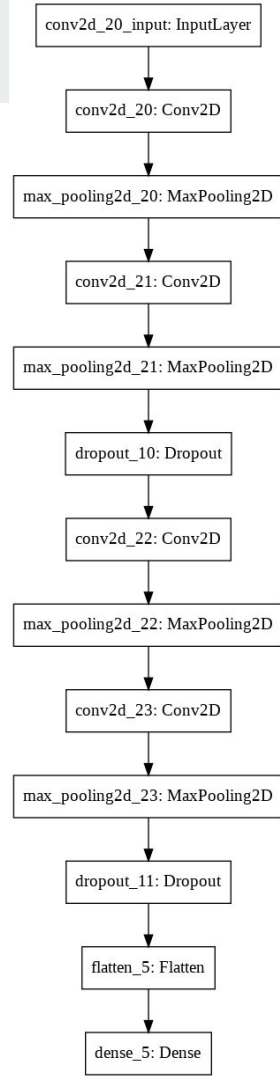
    model.add(Conv2D(64, (3, 3), activation='relu'))
    model.add(MaxPooling2D())

    model.add(Dropout(0.5))

    model.add(Flatten())



    # Binary classifier
    model.add(Dense(1, activation='sigmoid'))

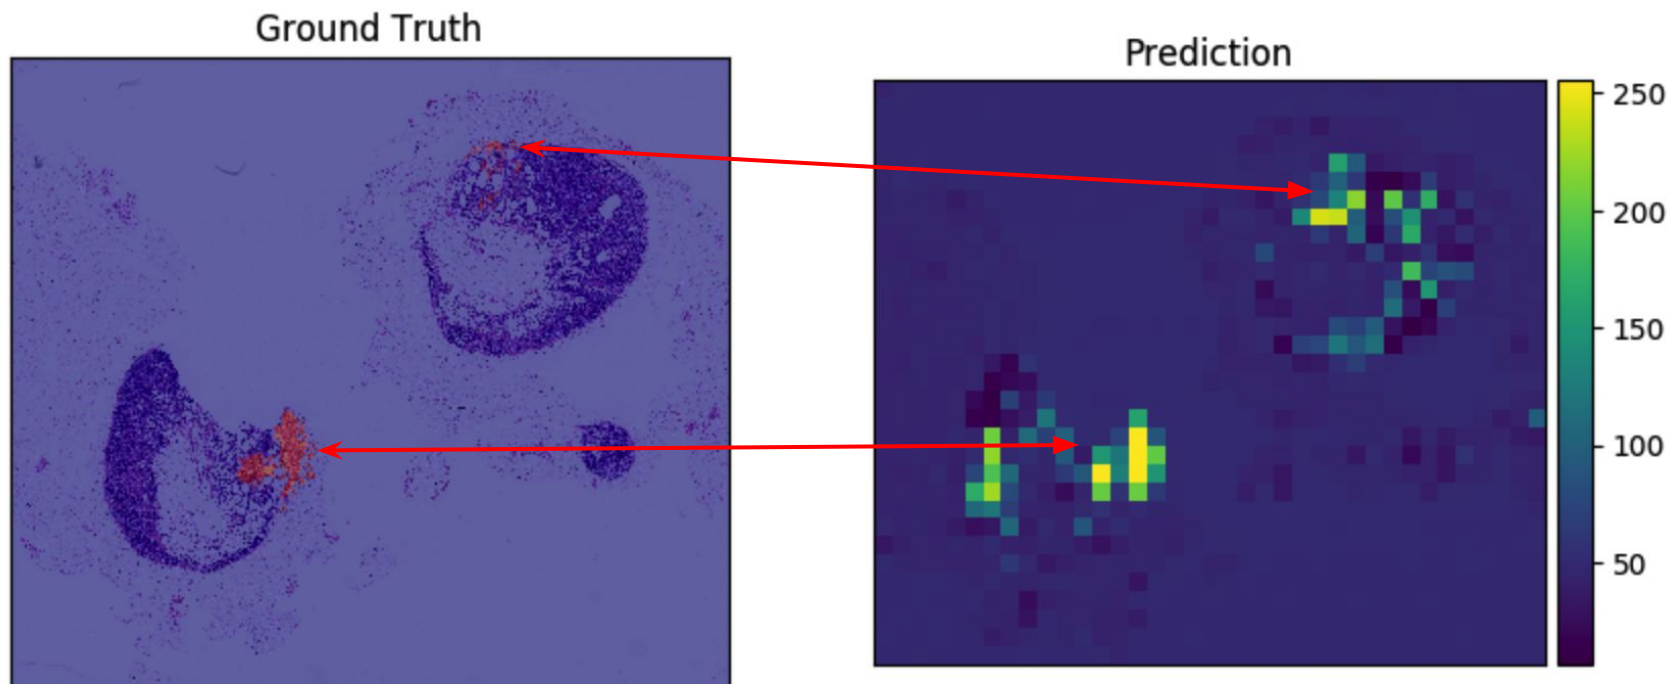
    return model
```



# ROC-AUC & Loss

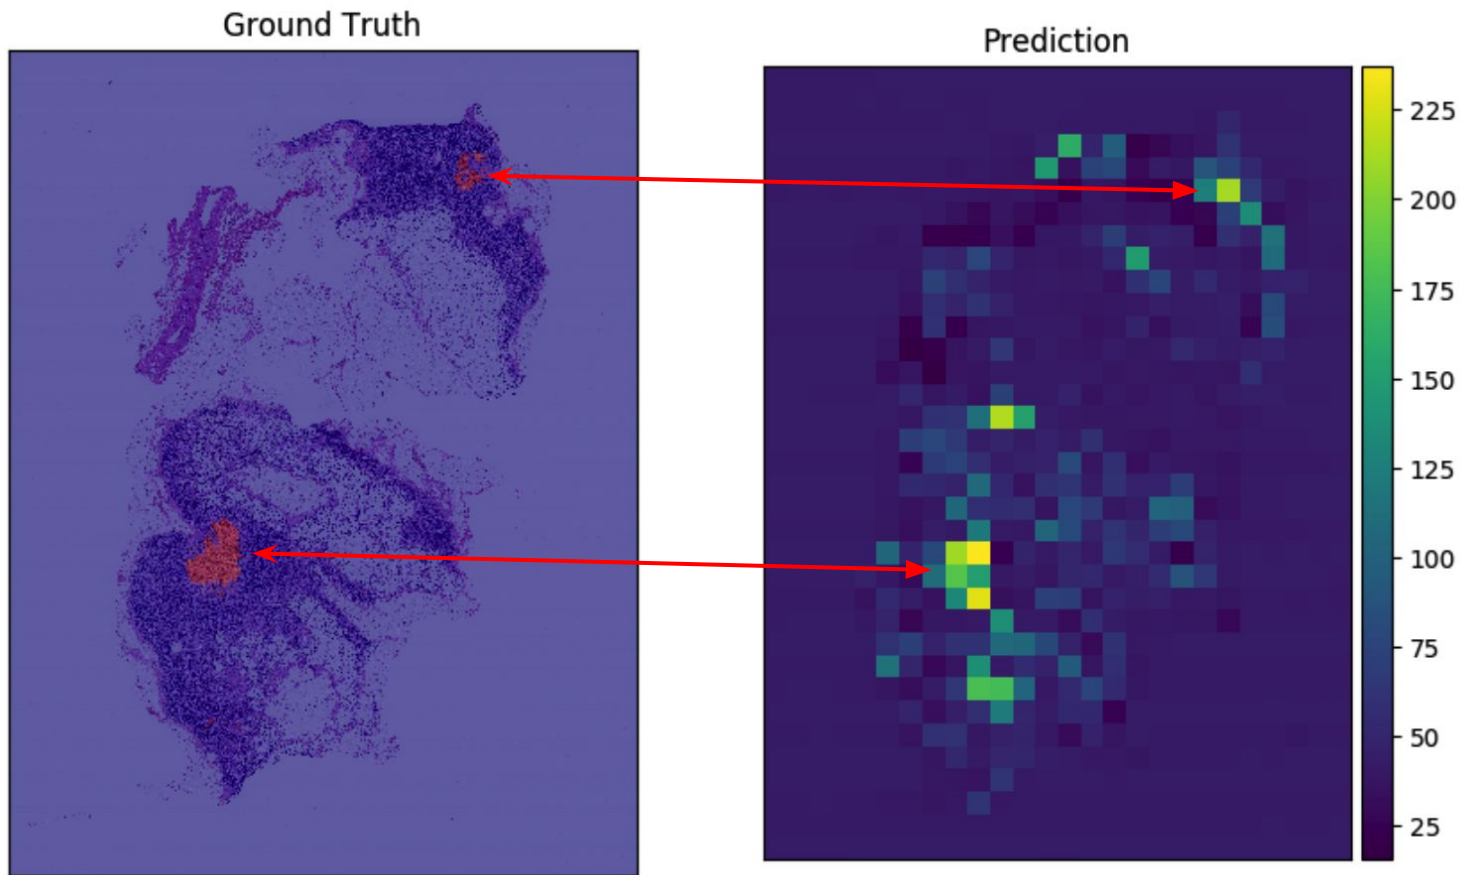
- ROC-AUC was used for imbalanced dataset
- Validation accuracy is higher/loss is lower sometimes
  - Regularization is applied during training but not validation
  - Training loss is measured during each epoch (thus is averaged over the epoch) while validation loss is measured after each epoch (thus is computed only once)
  - The validation set may be easier and there may also be leaks.
- Validation is decreasing at the last epoch
  - Not overfitting





Prediction: tumor\_075.tif





Prediction: tumor\_084.tif

# One-scale, InceptionV3

```
from tensorflow.keras.applications.inception_v3 import InceptionV3
from tensorflow.keras.layers import GlobalAveragePooling2D

def make_model_inception():
    model = Sequential()
    base_model = InceptionV3(include_top=False,
                             weights='imagenet',
                             input_shape=(299, 299, 3))

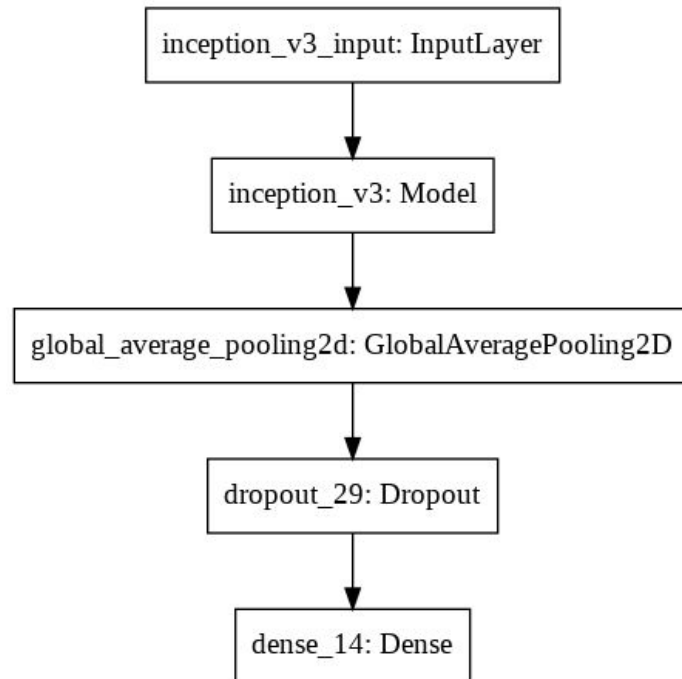
    # Do not update the pretrained weights during training
    base_model.trainable = False

    model.add(base_model)
    model.add(GlobalAveragePooling2D())

    # model.add(Dense(32, activation='relu'))
    model.add(Dropout(0.5))



    # Binary classifier
    model.add(Dense(1, activation='sigmoid'))

    return model
```




# ROC-AUC & Loss


- Epoch = 15
- Validation accuracy is always lower
- Validation increasing at epoch = 7
  - Suggesting early stop



Ground Truth




Prediction




Prediction: tumor\_075.tif

Ground Truth



Prediction



Prediction: tumor\_084.tif

# Multi-scale CNNs

```
from tensorflow.keras.layers import Input, concatenate

def make_model_multi():
    input1 = Input(shape=(299, 299, 3))
    input2 = Input(shape=(299, 299, 3))

    # The first branch operates on the first input: level 3
    x = Conv2D(32, (3, 3), activation='relu')(input1)
    x = MaxPooling2D()(x)
    x = Conv2D(32, (3, 3), activation='relu')(x)
    x = MaxPooling2D()(x)
    x = Dropout(0.5)(x)

    x = Conv2D(64, (3, 3), activation='relu')(x)
    x = MaxPooling2D()(x)
    x = Conv2D(64, (3, 3), activation='relu')(x)
    x = MaxPooling2D()(x)
    x = Dropout(0.5)(x)
    x = Flatten()(x)

    x = Model(inputs=input1, outputs=x)

    # The second branch operates on the second input: level 4
    y = Conv2D(32, (3, 3), activation='relu')(input2)
    y = MaxPooling2D()(y)
    y = Conv2D(32, (3, 3), activation='relu')(y)
    y = MaxPooling2D()(y)
    y = Dropout(0.5)(y)

    y = Conv2D(64, (3, 3), activation='relu')(y)
    y = MaxPooling2D()(y)
    y = Conv2D(64, (3, 3), activation='relu')(y)
    y = MaxPooling2D()(y)
    y = Dropout(0.5)(y)
    y = Flatten()(y)

    y = Model(inputs=input2, outputs=y)

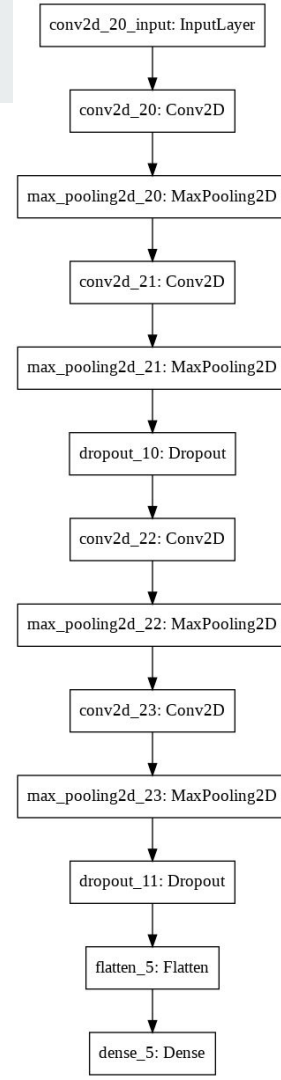
    # Combine the output of the two branches
    combined = concatenate([x.output, y.output])

    # Apply a FC layer and then a binary classifier
    z = Dense(2, activation='relu')(combined)

    z = Dense(1, activation='sigmoid')(z)



    model = Model(inputs=[x.input, y.input], outputs=z)

    return model
```




## ROC-AUC & Loss

- Epoch = 10
- Validation accuracy is higher/loss is lower at the start, but lower/loss is higher later
- Validation is decreasing at the last epoch
  - Not overfitting






Ground Truth




Prediction




Prediction: tumor\_075.tif



Ground Truth



Prediction



Prediction: tumor\_084.tif



# Conclusions

- The simple CNN has the best performance
  - InceptionV3 is the worst
  - Overfitting (small training set)
- Context patches are not so helpful
- Limitations
  - InceptionV3 fine tuning
  - Data augmentation
  - Larger scales, e.g., level 2